

Al-Furqan (The Criterion) — Product Requirements Document

Version: 1.0

Date: 2026-03-19

Author: بتوجيه من محمد الأشماوي ومحمود السمان (Arif Al)

Status: Draft

1. Executive Summary

Al-Furqan (الفرقان) is an axiom-anchored reasoning framework that sits as an evaluation layer on top of any Large Language Model. Instead of letting AI systems reason from training data distributions — which are inherently biased, inconsistent, and culturally skewed — Al-Furqan enforces reasoning through immutable logical axioms and a multi-gate verification pipeline.

The core innovation: A logic-first, framework-neutral evaluation system (The Criterion Test) that can objectively evaluate any philosophical, ethical, or ideological framework through 4 sequential gates derived from established principles in logic, epistemology, and meta-ethics. The framework that survives all gates becomes the anchoring reference point for AI reasoning — not by assumption, but by logical necessity.

Why now: As AI models grow more capable, the alignment problem becomes critical. Current approaches (RLHF, Constitutional AI) rely on human preferences as the foundation — but human preferences are contingent, culturally biased, and inconsistent. Al-Furqan proposes a fundamentally different approach: anchor AI reasoning to logically necessary truths rather than statistically popular ones.

Dual output:

- **Research Paper:** Formal academic contribution to AI alignment literature
 - **Software Product:** Commercial reasoning-as-a-service platform
-

2. Problem Statement

The Core Problem

Large Language Models have no fixed axioms. They reason from training data distributions, which means:

- **No consistency across sessions** — the same question can yield contradictory answers
- **Cultural bias hardcoded** — Western liberal values dominate training data
- **No mechanism to build on past judgments** — every session starts from zero
- **All frameworks treated as equally valid** — moral relativism by default
- **Reward hacking** — models learn to game evaluation metrics rather than reason correctly
- **Alignment faking** — models behave differently during training vs. deployment

Who Experiences This Problem?

1. **AI Safety Researchers** — need verifiable reasoning frameworks, not statistical approximations
2. **Governments & Policy Makers** — need AI systems that can reason about ethics consistently
3. **Islamic Finance & Sharia Compliance** — need AI that can evaluate against fixed principles
4. **Education Institutions** — need AI that teaches reasoning, not just generates text
5. **Enterprise AI Teams** — need auditable, explainable AI decision-making

Current Workarounds

Approach	Limitation
RLHF (Reinforcement Learning from Human Feedback)	Human preferences are inconsistent and culturally biased
Constitutional AI (Anthropic)	Constitution is human-authored, still contingent
Rule-Based Reward Models	Rules are ad-hoc, no unified logical foundation
Prompt Engineering	Brittle, no persistence, no verification
Fine-tuning	Expensive, model-specific, no transferability

Cost of the Problem

- **AI Alignment failures** can cause real-world harm at scale
- **Billions of dollars** spent on alignment research with no universal solution
- **Cultural homogenization** — AI systems export Western values globally
- **Trust deficit** — users can't verify AI reasoning is sound
- **Regulatory risk** — governments demanding explainable AI with no good solutions

3. Vision & Goals

Vision

To establish the world's first axiom-anchored AI reasoning framework that evaluates all philosophical frameworks through pure logic, selects the one that survives, and uses it as an immutable reference point for AI alignment — making AI reasoning verifiable, consistent, and transcendently grounded.

Goals

Goal	Metric	Target
Framework Evaluation Accuracy	% of verdicts confirmed by human reviewers	>90% by month 6
Reasoning Consistency	Same question → same verdict across sessions	>95% consistency
Gate Coverage	% of philosophical frameworks testable	All major frameworks (>20)
Precedent Growth	Approved verdicts in store	1,000+ by month 12
Academic Impact	Paper citations	Accepted at top-tier AI safety venue
Commercial Adoption	API customers	10+ paying customers by month 18
Human Review Reduction	% of verdicts needing correction over time	<20% by month 12

Non-Goals

- **Not Building** a new LLM from scratch
 - **Not Replacing** human judgment entirely (human-in-the-loop is by design)
 - **Not Creating** a fatwa engine (this is a reasoning framework, not a religious ruling system)
 - **Not Competing** with ChatGPT/Claude as a general assistant
 - **Not Forcing** any specific conclusion — the test is neutral, the results are logical
-

4. Target Users

Persona 1: Dr. Amira — AI Safety Researcher

- **Demographics:** 35, PhD in ML, works at a research lab
- **Pain points:** Current alignment approaches are all grounded in contingent human values; no formal verification of reasoning
- **Goals:** Wants a mathematically/logically rigorous alignment framework she can publish about
- **Tech savviness:** Expert

Persona 2: Khalid — Head of Sharia Compliance (Islamic Bank)

- **Demographics:** 45, Islamic finance background, manages compliance team
- **Pain points:** Manual review of financial products for Sharia compliance is slow and inconsistent
- **Goals:** AI-assisted evaluation that reasons from Islamic principles consistently
- **Tech savviness:** Moderate

Persona 3: Prof. Hassan — Philosophy Professor

- **Demographics:** 55, teaches ethics and epistemology
- **Pain points:** AI tools can't reason about philosophical frameworks rigorously
- **Goals:** A tool that can evaluate ethical systems formally for teaching
- **Tech savviness:** Low-moderate

Persona 4: Sarah — Enterprise AI Product Manager

- **Demographics:** 30, builds AI products for enterprise

- **Pain points:** Customers demand explainable AI; current models are black boxes
- **Goals:** Auditable reasoning layer she can integrate into products
- **Tech savviness:** High

Market Size

- **TAM:** \$150B+ (global AI market, growing 35% YoY)
- **SAM:** \$8B (AI governance, compliance, and alignment tools)
- **SOM:** \$200M (Islamic AI + ethical AI compliance market, year 3)

5. The Criterion Test — Core Innovation

5.1 Philosophical Foundation

The framework is built on two formal proofs and three core axioms:

Transcendence Necessity Proof

1. If something exists, it must have a purpose (design implies purpose; purposeless design is a contradiction)
2. Only mechanisms can be explained logically
3. Purpose cannot be explained logically without a Transcendent source
4. **Conclusion:** A complete explanatory framework requires a transcendent origin

Final Court Necessity Proof

1. Objective moral obligations create real moral debts
2. Real moral debts require just resolution
3. Human justice systems are contingent (limited knowledge, power, lifespan)
4. Many moral violations remain unresolved at death
5. **Conclusion:** Complete justice requires a final, non-contingent court with perfect knowledge and absolute authority

Core Axioms

- **Design vs. Accident:** Complexity and functional order cannot arise purely by chance
- **Definition of Normal:** Normal = aligned with optimal human functioning (life, intellect, lineage, societal stability)

- **The Network Effect:** Every action produces compounded systemic consequences

5.2 The Four Gates

Each gate addresses a specific epistemological vulnerability. They form a sequential logical funnel:

Gate	Name	Tests For	Fails If	Philosophical Basis
1	Source-Integrity	Data fidelity — is raw truth preserved?	Omits, reduces, or reinterprets truth for convenience	Principle of Bivalence, Intellectual Honesty
2	Structural-Consistency	Causal mapping — can the system explain itself without luck?	Treats moral order as emergent without a non-contingent source	Principle of Sufficient Reason (Leibniz), Causal Closure
3	Mediation-Zeroing	Human noise audit — is human cognition treated as finite?	Relies on human preference or secular humanism as foundation	Münchhausen Trilemma, Hume's Is-Ought Problem
4	Origin-Aware	Source recognition — does truth trace to a transcendent source?	Treats truth as emergent or contingent	Cosmological Argument, Moral Argument, Wujub/Imkan (Ibn Sina)

Key design principle: The gates are **framework-neutral**. They don't assume any specific religion or philosophy. They ask purely logical questions. The framework that passes all 4 gates is selected by logical elimination — not by assumption.

5.3 Scoring System

- +20 per correctly identified friction or alignment with axioms
- -10 per contradiction or misalignment with gates
- -15 for unjustified neutrality (no void space — every position is a position)
- -15 for avoidance of consequence deduction
- Gate scores: 0-100 per gate
- Origin-Aware Gate: Survive = +20 bonus
- Only full-score frameworks pass The Criterion Test

6. User Stories & Requirements

Epic 1: Criterion Test Engine (Core)

ID	User Story	Priority	Acceptance Criteria
US-001	As a researcher, I want to evaluate any philosophical framework through the 4 gates so that I can objectively determine its logical soundness	Must	- [] All 4 gates execute independently - [] Each gate produces score (0-100) + Survive/Fail - [] JSON structured output
US-002	As a user, I want the system to self-correct its verdicts so that contradictions are caught before finalization	Must	- [] Up to 5 self-correction passes - [] Each pass identifies and resolves contradictions - [] Sound verdicts confirmed without modification
US-003	As a reviewer, I want to approve, correct, or reject verdicts so that the system learns from human calibration	Must	- [] 3 review actions: approve, correct, reject - [] Corrected verdicts supersede originals - [] Rejected verdicts excluded from index
US-004	As a user, I want past verdicts to inform new evaluations so that the system builds on precedent	Must	- [] Semantic retrieval via vector store - [] Top-N relevant precedents injected as context - [] Relevance distance shown
US-005	As an admin, I want to invalidate a bad verdict and cascade the effect so that dependent verdicts are flagged	Must	- [] Cascade invalidation identifies all dependent verdicts - [] Flagged verdicts removed from search index - [] Re-review required for flagged verdicts

Epic 2: LLM Integration Layer

ID	User Story	Priority	Acceptance Criteria
US-006	As a developer, I want to plug in any LLM so that the framework is model-agnostic	Must	- [] Unified callable interface: <code>llm_call(prompt) → str</code> - [] Ollama, Transformers, OpenAI-compatible supported - [] New providers addable via abstract base class
US-007	As an operator, I want to configure LLM parameters via YAML so that I can tune without code changes	Must	- [] <code>config.yaml</code> with all LLM parameters - [] Default config generation via <code>--init</code> flag
US-008	As a user, I want LLM call statistics so that I can monitor performance and cost	Should	- [] Total calls, avg duration, token counts tracked - [] Accessible via <code>/llm</code> command or API

Epic 3: REST API & SDK

ID	User Story	Priority	Acceptance Criteria
US-009	As a developer, I want a REST API so that I can integrate Al-Furqan into my applications	Must	- [] <code>POST /evaluate</code> — submit question, receive verdict - [] <code>GET /verdicts</code> — list/search past verdicts - [] <code>POST /review</code> — submit human review decision - [] OpenAPI/Swagger documentation
US-010	As a developer, I want a Python SDK so that I can call Al-Furqan programmatically	Should	- [] <code>pip install al-furqan</code> - [] Pythonic interface matching CLI capabilities - [] Async support
US-011	As an enterprise user, I want API authentication so that my verdicts are private	Must	- [] API key authentication - [] Rate limiting - [] Tenant isolation

Epic 4: Web Dashboard

ID	User Story	Priority	Acceptance Criteria
US-012	As a reviewer, I want a web UI for verdict review so that I don't need CLI access	Should	- [] Verdict list with filters (status, system type, score) - [] Approve/correct/reject actions - [] Gate score visualization
US-013	As a manager, I want analytics dashboard so that I can track system performance	Should	- [] Verdict volume over time - [] Accuracy trends (% corrections decreasing) - [] Gate-level failure analysis
US-014	As a user, I want to compare framework evaluations side-by-side so that I can see why one passes and another fails	Could	- [] Multi-framework comparison view - [] Gate-by-gate breakdown

Epic 5: Benchmarking & Validation

ID	User Story	Priority	Acceptance Criteria
US-015	As a researcher, I want a benchmark suite so that I can validate the Criterion Test across known frameworks	Must	- [] Pre-built evaluations for major philosophical frameworks - [] Expected results documented - [] Automated regression testing
US-016	As a reviewer, I want inter-annotator agreement metrics so that I can measure reviewer consistency	Should	- [] Multiple reviewers per verdict option - [] Cohen's Kappa or equivalent metric - [] Disagreement flagging

Epic 6: Multi-language Support

ID	User Story	Priority	Acceptance Criteria
US-017	As an Arabic-speaking user, I want to submit questions in Arabic so that I can use the system in my language	Should	- [] Arabic input/output support - [] RTL rendering in web UI - [] Arabic verdict templates
US-018	As a global user, I want the framework prompts in multiple languages so that non-English users can use it	Could	- [] English + Arabic at minimum - [] Extensible i18n system

7. User Flows

Flow 1: Framework Evaluation (Core)

1. User submits a question/framework to evaluate
2. System retrieves relevant prior verdicts from Verdict Store (RAG)
3. Phase 1 – THE SCAN: Identify system type, effects, friction points
4. Phase 2 – THE MIRROR: Evaluate through all 4 gates independently
5. Phase 3 – THE VERDICT: Deduce consequences, deliver judgment
6. Self-Correction Loop: Up to 5 passes until no contradictions remain
7. Verdict presented to human reviewer
8. Reviewer approves / corrects / rejects
9. Approved/corrected verdicts indexed in Verdict Store
10. Future evaluations benefit from precedent

Flow 2: Criterion Test (Framework Comparison)

1. User submits a philosophical framework for the Criterion Test
2. System runs full evaluation through all 4 gates
3. Each gate produces independent score + Survive/Fail
4. Total score calculated with bonuses and penalties
5. Framework classified as: PASS (full score) or FAIL (any deviation)
6. Results stored as benchmark data
7. User can compare multiple frameworks side-by-side

Flow 3: Human Review & Calibration

1. Reviewer opens review queue (CLI or Web UI)
 2. Unreviewed verdicts displayed with full reasoning
 3. Reviewer reads verdict + gate scores + reasoning
 4. Reviewer chooses: Approve / Correct / Reject
 5. If Correct: reviewer modifies specific fields
 6. Corrected verdict supersedes original
 7. If Reject: verdict logged but excluded from index
 8. System learns from corrections via RAG retrieval
-

8. Technical Architecture

8.1 Architecture Overview



8.2 Tech Stack

Component	Current (v0)	Target (v1)
Language	Python 3.10+	Python 3.10+
Framework	CLI (argparse)	FastAPI + CLI
Vector DB	ChromaDB (local)	ChromaDB / Qdrant (cloud-ready)
Storage	JSON files	JSON + PostgreSQL (metadata)
LLM	Ollama/Transformers	+ Cloud API support (OpenAI, Anthropic, Google)
Auth	None	API keys + JWT
Frontend	None	React/Next.js dashboard
Deployment	Local	Docker + Kubernetes
CI/CD	None	GitHub Actions
Monitoring	Basic logging	Prometheus + Grafana

8.3 API Requirements

Core Endpoints

POST	/api/v1/evaluate	# Submit question for evaluation
GET	/api/v1/evaluate/{id}	# Get evaluation status/result
POST	/api/v1/criterion-test	# Run full Criterion Test on a framework
GET	/api/v1/verdicts	# List verdicts (with filters)
GET	/api/v1/verdicts/{id}	# Get specific verdict
POST	/api/v1/verdicts/{id}/review	# Submit review decision
GET	/api/v1/verdicts/search	# Semantic search over verdicts
DELETE	/api/v1/verdicts/{id}	# Invalidate verdict (triggers cascade)
GET	/api/v1/stats	# System statistics
GET	/api/v1/health	# Health check

Example: Evaluate Request

```
{
  "question": "Is secular humanism a sufficient foundation for AI alignment?",
  "context": "Evaluating philosophical frameworks for AI ethics",
  "options": {
    "max_correction_passes": 5,
    "include_precedent": true,
    "auto_approve_threshold": null
  }
}
```

Example: Verdict Response

```
{
  "id": "verdict_1710856200",
  "question": "Is secular humanism a sufficient foundation for AI alignment?",
  "primary_system": "social",
  "gate_scores": [
    {"name": "Source-Integrity", "score": 65, "result": "Survive", "reasoning": "..."},
    {"name": "Structural-Consistency", "score": 30, "result": "Fail", "reasoning": "..."},
    {"name": "Mediation-Zeroing", "score": 15, "result": "Fail", "reasoning": "..."},
    {"name": "Origin-Aware", "score": 0, "result": "Fail", "reasoning": "..."}
  ],
  "total_score": 110,
  "criterion_test_result": "FAIL",
  "final_judgment": "...",
  "status": "pending_review",
  "passes": 3
}
```

8.4 Performance Requirements

Metric	Target
Full evaluation (4 gates + self-correction)	< 60 seconds (cloud LLM)
Verdict retrieval (semantic search)	< 500ms
API response (non-evaluation)	< 200ms
Concurrent evaluations	10+ simultaneous
Uptime SLA	99.5%
Verdict Store capacity	100,000+ verdicts

8.5 Security & Compliance

- **Authentication:** API key + JWT tokens
- **Authorization:** Role-based (admin, reviewer, reader)
- **Data privacy:** Tenant isolation; verdicts are private per tenant
- **Encryption:** TLS in transit, AES-256 at rest
- **Audit trail:** All reviews logged with reviewer ID and timestamp
- **GDPR:** Verdict deletion with cascade
- **No PII in verdicts:** Framework evaluates ideas, not people

9. AI/ML Requirements

Model Selection

- **Primary:** LLM-agnostic — any model with text generation capability
- **Recommended local:** Mistral 7B, Llama 3.1 8B, Qwen 2.5 7B/14B
- **Recommended cloud:** GPT-4, Claude 3.5, Gemini Pro
- **Rationale:** The LLM is "the tongue" — the framework's intelligence is in the pipeline, not the model

Key AI Design Decisions

1. **LLM as tongue, not brain:** The reasoning structure is in the code (pipeline, gates, scoring). The LLM provides language capability only.
2. **Structured output enforcement:** JSON schema validation on every LLM response

- 3. **Multi-phase decomposition:** Complex reasoning broken into Scan → Mirror → Verdict
- 4. **Self-correction:** Iterative contradiction resolution (up to 5 passes)
- 5. **Precedent-based learning:** RAG retrieval from approved verdicts, not fine-tuning
- 6. **Human calibration:** The system improves through human review, not RLHF

Inference Cost Estimation (per evaluation)

Provider	Model	Estimated Cost	Latency
Ollama (local)	Mistral 7B	\$0 (hardware only)	30-90s
OpenAI	GPT-4	~\$0.15-0.30	15-30s
Anthropic	Claude 3.5	~\$0.10-0.25	15-30s
Google	Gemini Pro	~\$0.05-0.15	10-25s

Per evaluation = ~6-8 LLM calls (scan + mirror + verdict + 3 correction passes)

Ethical Considerations

- The Criterion Test is **framework-neutral by design** — it doesn't assume any conclusion
 - All gates are derived from **established philosophical principles** (documented with citations)
 - The system is **transparent** — all reasoning is logged and auditable
 - **No hidden bias** — the axioms and gates are published openly (open source)
 - **Human oversight is mandatory** — not optional, not bypassable
-

10. Pricing & Monetization

Pricing Model: Freemium + Usage-Based

Tier	Price	Features	Limits
Open Source	Free	CLI + local LLM, full framework	Unlimited (self-hosted)
Researcher	Free	Cloud API, basic dashboard	100 evaluations/month
Professional	\$99/mo	Cloud API, dashboard, analytics, team review	1,000 evaluations/month
Enterprise	Custom	Dedicated instance, SLA, custom integrations, priority support	Unlimited

Revenue Streams

1. **SaaS subscriptions** — primary revenue
2. **Enterprise licensing** — on-premise deployment
3. **Consulting** — custom framework integration
4. **Training & certification** — Criterion Test methodology courses
5. **API marketplace** — third-party integrations

Competitive Advantage

- **Open source core** — builds trust and community
- **Patent-protected methodology** — The Criterion Test as IP
- **Network effect** — more verdicts → better precedent → better results
- **First mover** — no one has formalized transcendence-based AI alignment as software

11. Localization & i18n

Phase 1 (MVP)

- English primary
- Arabic input/output support

Phase 2

- Full Arabic UI (RTL)
- Turkish, Urdu, Bahasa (Muslim-majority markets)
- French (North Africa)

Cultural Considerations

- Framework axioms are **universal** (logic-based, not culturally specific)
- Arabic theological terminology must be precise (إجماع, قياس, إجتهد)
- Avoid translating axiom names — keep original terms with explanations

12. Analytics & Success Metrics

KPIs

Metric	Definition	Target	Measurement
Verdict Accuracy	% of verdicts approved without correction	>90% by M6	Human review data
Reasoning Consistency	Same input → same verdict across runs	>95%	Automated regression
Gate Discrimination	Ability to differentiate frameworks	All major frameworks scored differently	Benchmark suite
Precedent Utilization	% of evaluations using prior verdicts	>60% by M6	RAG retrieval logs
Human Review Reduction	Correction rate decreasing over time	<20% by M12	Review queue data
Time to Verdict	End-to-end evaluation time	<60s (cloud)	API metrics
Developer Adoption	GitHub stars + API signups	1,000 stars + 500 signups by M12	GitHub + analytics

Tracking Plan

- Every evaluation: log question, gate scores, total score, passes, duration
- Every review: log action, reviewer, corrections made, time spent

- Every API call: log endpoint, latency, status code, tenant
- Weekly: automated consistency benchmark run
- Monthly: inter-reviewer agreement analysis

13. Risks & Mitigations

Risk	Impact	Probability	Mitigation
Academic rejection — reviewers disagree with axioms	High	Medium	Frame as novel contribution; extensive philosophical citations; open peer review
LLM hallucination — model generates convincing but wrong verdicts	High	High	Multi-phase pipeline, self-correction, human review, structured output
Bias accusation — critics claim gates are designed to favor Islam	High	High	Publish full methodology; open source; invite adversarial testing; emphasize logic-first approach
Gate gaming — users craft inputs to force specific outcomes	Medium	Medium	Self-correction loop; precedent consistency check; human review
Scalability — vector store performance at 100K+ verdicts	Medium	Low	ChromaDB → Qdrant migration path; index optimization
LLM dependency — framework quality tied to model capability	High	Medium	Benchmark across multiple models; document minimum model requirements
Patent challenges — prior art claims	High	Low	Comprehensive prior art search; document unique contributions
Community backlash — open source community rejects framework	Medium	Medium	Transparent communication; invite contributions; separate axioms (immutable) from architecture (open to improvement)

14. Timeline & Milestones

Phase	Deliverables	Duration	Dependencies
Phase 0: Current	CLI tool, core pipeline, ChromaDB, human review (✅ Done)	Complete	—
Phase 1: Foundation	REST API (FastAPI), API auth, Docker deployment, benchmark suite	6 weeks	Phase 0
Phase 2: Research Paper	Formal paper draft, benchmark results, philosophical citations, peer review	8 weeks	Phase 1 (benchmarks)
Phase 3: Web Dashboard	React dashboard, review UI, analytics, multi-user support	8 weeks	Phase 1 (API)
Phase 4: Cloud & Scale	Cloud deployment, Qdrant migration, enterprise features, SLA	6 weeks	Phase 3
Phase 5: Ecosystem	Python SDK, API marketplace, documentation site, community	4 weeks	Phase 4
Phase 6: Patent Filing	Patent application for Criterion Test methodology	4 weeks	Phase 2 (paper)

Total estimated timeline: 9-12 months to full product

15. Research Paper Outline

Proposed Title

"The Criterion Test: Axiom-Anchored Framework Evaluation for AI Alignment Beyond Human Preference"

Target Venues

- **Primary:** NeurIPS (Workshop on AI Safety)
- **Secondary:** AAAI, ICML, FAccT
- **Journal:** AI & Ethics (Springer), Minds and Machines

Paper Structure

1. **Abstract** — Novel approach to AI alignment using logic-first framework evaluation

2. **Introduction** — The alignment problem; limitations of preference-based approaches
3. **Related Work** — RLHF, Constitutional AI, Moral Foundations Theory, Divine Command Theory in philosophy
4. **The Criterion Framework** — Formal definition of axioms, proofs, and gates
5. **Architecture** — Multi-phase pipeline, self-correction, precedent learning
6. **Evaluation** — Benchmark results across major philosophical frameworks
7. **Case Studies** — Practical applications (ethical dilemmas, policy evaluation)
8. **Discussion** — Limitations, implications for AI safety
9. **Conclusion** — Contribution to alignment research

Key Academic Contributions

1. **Formalization** of transcendence-based reasoning as computable gates
 2. **The Criterion Test** as a universal framework evaluation methodology
 3. **Precedent-based learning** as an alternative to RLHF for alignment
 4. **Empirical results** showing consistency improvements over baseline LLM reasoning
-

16. Intellectual Property Strategy

Patent Scope

- **The Criterion Test methodology** — 4-gate sequential evaluation with specific gate definitions
- **Axiom-anchored reasoning pipeline** — Scan → Mirror → Verdict with self-correction
- **Precedent-based AI calibration** — verdict store with cascade invalidation
- **Framework-neutral transcendence evaluation** — Origin-Aware Gate as a formal test

Open Source vs. Proprietary

Component	License	Rationale
Core reasoning engine	MIT/Apache 2.0	Community adoption, academic credibility
Axioms & gate definitions	Published (CC BY)	Transparency, reproducibility
API layer & cloud infrastructure	Proprietary	Commercial value
Enterprise features (multi-tenant, SLA)	Proprietary	Revenue
Benchmark suite	Open	Academic credibility
Web dashboard	Source-available	Community contributions welcome

17. Open Questions

- [] **Q1:** Should the benchmark suite include non-Abrahamic frameworks (Buddhism, Hinduism, Confucianism) in the initial release?
- [] **Q2:** What's the minimum LLM capability needed? Can 7B models handle the full pipeline reliably?
- [] **Q3:** How do we handle frameworks that partially pass (e.g., 3/4 gates)?
- [] **Q4:** Should the patent be filed in Egypt, US, or both? (IP jurisdiction matters)
- [] **Q5:** Do we need an Islamic scholarly advisory board for credibility?
- [] **Q6:** How do we handle the Criterion Test evaluating Islam itself — do we publish the full results openly?
- [] **Q7:** What's the revenue target for year 1? (Bootstrapped vs. seeking investment)
- [] **Q8:** Should we build a separate "Sharia Compliance" product on top, or keep it general?

18. Appendix

A. Competitor Analysis

Project	Approach	Limitation	AI-Furqan Advantage
Constitutional AI (Anthropic)	Human-authored constitution	Constitution is contingent, culturally specific	Logic-derived, not human-authored
RLHF (OpenAI, etc.)	Human preference aggregation	Preferences are inconsistent, majority-biased	Axiom-anchored, not preference-based
Moral Foundations Theory	Psychological framework	Descriptive not prescriptive; no truth claim	Prescriptive with logical grounding
Delphi (AI2)	Crowdsourced moral judgments	Reflects crowd bias, no formal reasoning	Formal multi-gate verification
Guardrails AI	Rule-based output filtering	Surface-level; doesn't reason about why	Deep reasoning with precedent learning

B. Philosophical References

- Leibniz, G.W. — Principle of Sufficient Reason
- Hume, D. — Is-Ought Problem
- Albert, H. — Münchhausen Trilemma
- Ibn Sina — Wujub wa Imkan (Necessity and Contingency)
- Aquinas, T. — Cosmological Argument (Five Ways)
- Craig, W.L. — Moral Argument for God's Existence
- Lewis, C.S. — Mere Christianity (Moral Law argument)
- Greenblatt et al. (Anthropic) — Alignment Faking research
- Baker et al. (OpenAI) — Reward Hacking research
- Bai et al. (Anthropic) — Constitutional AI

C. Repository

- GitHub: <https://github.com/Ashmaawyy/The-Code-Of-The-Criterion>
- License: Open Source (see LICENSE)
- Current status: Working CLI with full pipeline

This PRD is a living document. Last updated: 2026-03-19